

GURU JAMBHESHWAR
UNIVERSITY OF SCIENCE AND
TECHNOLOGY, HISAR



MINOR PROJECT
SESSION – 2023-24

SUBMITTED BY :-
AAKASH SINHA
210010159004
B.TECH. CSE-AI&ML

SUBMITTED TO :-
CSE Dept.
GJUS&T, HISAR

DECLARATION

I hereby declare that work presented in work entitled "**Iris Flower prediction**" as Minor Project Assignment is solely made by me with guidance of Data science and Machine Learning Course online. And take training from Robatic club(IIT Guwahati) I assure that this project is solely my hard work and not copied from anywhere.

Aakash Sinha
(210010159004)
B.Tech. CSE-AI&ML

ACKNOWLEDGEMENT

The success and final outcome of this project required a lot of guidance and assistance from many people and I am extremely fortunate to have got this all along the completion of our project. Whatever I have done is only due to such guidance and assistance and we would not forget to thank them.

I express my deepest gratitude to all the teachers of my university for extending their inevitable and valuable support to us.

I am thankful to and fortunate enough to get constant encouragement, support and guidance from our faculty members without whom this project would have been a distant reality. I would also like to extend my sincere regards to my family and well-wishers for their timely support.

CONTENTS

- Certificate
- Introduction about the Project
- Tech Stack Used
- Source Code
- Machine Learning Algorithm
- Conclusion

❖ Certificate of Achievement:-



INTRODUCTION ABOUT THE PROJECT

The Iris Flower Classification task involves the categorization of iris flowers into three species: Setosa, Versicolor, and Virginica, based on certain features.

The dataset used for this classification is the famous Iris dataset, containing measurements of sepal length, sepal width, petal length, and petal width.

2. Dataset Overview:-

The Iris dataset consists of 150 samples, with 50 samples from each of the three species.

Each sample has four features: sepal length, sepal width, petal length, and petal width.

The dataset is well-balanced, making it suitable for training and evaluating machine learning models.

3. Data Preprocessing:-

Checked for missing values: Ensure that there are no missing values in the dataset.

Exploratory Data Analysis (EDA): Visualized the distribution of each feature to understand the data.

Feature Scaling: Standardized or normalized the features to bring them to a similar scale.

4. Model Selection:

Chose appropriate machine learning algorithms for classification, such as:

Logistic Regression

Decision Trees

Support Vector Machines (SVM)

k-Nearest Neighbors (k-NN)

Random Forest

Split the dataset into training and testing sets for model evaluation.

5. Model Training:

Trained the selected models using the training dataset.

Utilized techniques like cross-validation to ensure robust model performance.

Tuned hyperparameters to enhance model accuracy.

6. Model Evaluation:

Evaluated the models on the testing set using metrics such as accuracy, precision, recall, and F1 score.

Visualized the results using confusion matrices and ROC curves.

7. Results:

Summarized the performance of each model and selected the best-performing one based on evaluation metrics.

Discussed any insights gained from feature importance analysis.

TECH Stack USED

When creating a web application for Iris flower classification using Streamlit, you can leverage a specific technology stack to ensure a smooth development process. Below is a suggested tech stack:

Streamlit:-

Streamlit is chosen for its simplicity and ease of use in creating interactive web applications with Python. It allows for the rapid development of data applications without the need for extensive web development knowledge.

Python:-

Python is the programming language of choice due to its wide adoption in the data science and machine learning community. Streamlit is a Python library, making it a natural choice for this project.

Pandas:-

Pandas is used for data manipulation and analysis. It is instrumental in handling the Iris dataset, allowing for easy loading, exploration, and manipulation of data before feeding it into machine learning models.

Scikit-Learn:-

Scikit-Learn provides a comprehensive set of tools for machine learning in Python. It includes various algorithms, preprocessing

techniques, and evaluation metrics. For Iris flower classification, Scikit-Learn's RandomForestClassifier is used, but other classifiers can also be explored.

Joblib:-

Joblib is utilized for model persistence, allowing the trained machine learning model to be saved and loaded efficiently. This is important in a web application setting where you might want to reuse a pre-trained model without retraining it every time the app is launched.

NumPy:-

NumPy is a foundational library for numerical operations in Python. While not explicitly used in the Streamlit app, it often underlies operations within Pandas and Scikit-Learn.

Matplotlib and Seaborn:-

Matplotlib and Seaborn are used for data visualization. In a Streamlit app, these libraries can help create insightful plots and graphs to showcase the Iris dataset and visualize model performance metrics.

Virtual Environment:-

A virtual environment is created to manage the project's dependencies, ensuring that the required libraries and their versions are isolated from other projects. This helps in maintaining consistency and avoiding conflicts between different projects.

Version Control (Optional):

Version control tools like Git can be used to track changes in the codebase, collaborate with others, and revert to previous versions if needed.

IDE (Integrated Development Environment):-

An Integrated Development Environment such as Visual Studio Code, PyCharm, or Jupyter Notebooks is chosen for efficient coding, debugging, and collaboration.

❖ **Machine Learning Algorithm:-**

Machine learning algorithms are pieces of code that help people explore, analyze, and find meaning in complex data sets. Each algorithm is a finite set of unambiguous step-by-step instructions that a machine can follow to achieve a certain goal.

Here are many machine Learning Algorithm for train model, such as :

- Linear regression
- Logistic regression
- Decision tree
- SVM algorithm
- Naive Bayes algorithm
- KNN algorithm
- K-means
- Random forest algorithm
- Dimensionality reduction algorithms
- Gradient boosting algorithm and AdaBoosting algorithm

❖ Random Forest Classifier:-

Random forest is one of the most popular tree-based supervised learning algorithms. It is also the most flexible and easy to use.

The algorithm can be used to solve both classification and regression problems. Random forest tends to combine hundreds of **decision trees** and then trains each decision tree on a different sample of the observations.

The final predictions of the random forest are made by averaging the predictions of each individual tree.

The benefits of random forests are numerous. The individual decision trees tend to **overfit** to the training data but random forest can mitigate that issue by **averaging** the prediction results from different trees. This gives random forests a higher predictive accuracy than a single decision tree.

The random forest algorithm can also help you to find features that are **important** in your dataset. It lies at the base of the [Boruta algorithm](#), which selects important features in a dataset.

Random forest has been used in a variety of applications, for example to provide recommendations of different products to customers in e-commerce.

In medicine, a random forest algorithm can be used to identify the patient's disease by analyzing the patient's medical record.

Also in the banking sector, it can be used to easily determine whether the customer is fraudulent or legitimate.

How does the Random Forest algorithm work?

The random forest algorithm works by completing the following steps:


Step 1: The algorithm select random samples from the dataset provided.

Step 2: The algorithm will create a decision tree for each sample selected. Then it will get a prediction result from each decision tree created.

Step 3: Voting will then be performed for every predicted result. For a classification problem, it will use **mode**, and for a regression problem, it will use **mean**.

Step 4: And finally, the algorithm will select the most voted prediction result as the final prediction.

Random Forest Classifier



how it works

Random Forest in Practice

Now that you know the ins and outs of the random forest algorithm, let's build a random forest classifier.

SOURCE CODE

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 from sklearn.metrics import accuracy_score
        5 from sklearn.ensemble import RandomForestClassifier
        6
```

```
In [2]: 1 # load iris dataset
        2
        3 data = pd.read_csv("Iris.csv")
        4 data
```

```
Out[2]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...	...	...	...	...	...	...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows x 6 columns

```
In [3]: 1 # delete the column 'Id'
        2
        3 df = data.drop(columns = ['Id'])
        4 df.head()
```

```
Out[3]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [4]: 1 # check any missing values
        2
        3 df.isna()
        4
```

```
Out[4]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...	...	...	...	...	...
145	False	False	False	False	False
146	False	False	False	False	False
147	False	False	False	False	False
148	False	False	False	False	False

```
In [5]: 1 # display a stats about this data
2
3 df.describe()
```

```
Out[5]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
In [6]: 1 # information about this data
2
3 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   SepalLengthCm         150 non-null   float64
1   SepalWidthCm          150 non-null   float64
2   PetalLengthCm         150 non-null   float64
3   PetalWidthCm          150 non-null   float64
4   Species                150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [7]: 1 # display no. of samples on each class
2 df['Species'].value_counts()
```

```
Out[7]: Iris-setosa      50
Iris-versicolor      50
Iris-virginica       50
Name: Species, dtype: int64
```

preprocessing the dataset

```
In [8]: 1 # check any missing values
2
3 df.isna()
4
```

```
Out[8]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...	...	...	...	...	...
145	False	False	False	False	False
146	False	False	False	False	False
147	False	False	False	False	False
148	False	False	False	False	False
149	False	False	False	False	False

150 rows x 5 columns

```

In [9]: # CHECK FOR NULL VALUES
2
3 df.isnull().sum()

```

```

Out[9]: SepalLengthCm    0
        SepalWidthCm     0
        PetalLengthCm    0
        PetalWidthCm     0
        Species         0
        dtype: int64

```

Exploratory Data Analysis

```





In [10]: 1 # Histogram Plot
         2
         3 df.hist(bins=None, figsize=(12,8))

```

```

Out[10]: array([[<Axes: title={'center': 'SepalLengthCm'}>,
                <Axes: title={'center': 'SepalWidthCm'}>],
               [<Axes: title={'center': 'PetalLengthCm'}>,
                <Axes: title={'center': 'PetalWidthCm'}>]], dtype=object)

```



use seaborn


```

In [11]: 1 import seaborn as sns
         2 import matplotlib.pyplot as plt
         3
         4
         5 # Use seaborn to create pair plots for all features
         6 sns.set(style="ticks")
         7 sns.pairplot(df, hue="Species", markers=["o", "s", "D"])
         8
         9
        10

```

8
9
10

Out [11]: <seaborn.axisgrid.PairGrid at 0x1684a1250>



In []: 1

correlation matrix

In [12]: 1
2 df.corr()

/var/folders/vz/qkwbtk057c5bs6dyj_w4d8f80000gn/T/ipykernel_6381/143363973.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
df.corr()

Out [12]:

	Sepal.LengthCm	Sepal.WidthCm	Petal.LengthCm	Petal.WidthCm
Sepal.LengthCm	1.000000	-0.109369	0.871754	0.817954
Sepal.WidthCm	-0.109369	1.000000	-0.420516	-0.356544
Petal.LengthCm	0.871754	-0.420516	1.000000	0.962757
Petal.WidthCm	0.817954	-0.356544	0.962757	1.000000


In [13]: 1 corr = df.corr()
2 fig,ax = plt.subplots(figsize=(10,10))
3 sns.heatmap(corr,annot = True, ax = ax, cmap = 'coolwarm')

/var/folders/vz/qkwbtk057c5bs6dyj_w4d8f80000gn/T/ipykernel_6381/3904265891.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
corr = df.corr()

Out [13]: <Axes: >

2 fig,ax = plt.subplots(figsize=(6,6))
3 sns.heatmap(corr,annot = True, ax = ax, cmap = 'coolwarm')

Out [32]: <Axes: >



Label Encoder

```
In [14]: 1 from sklearn.preprocessing import LabelEncoder
        2 le = LabelEncoder()
```

```
In [15]: 1 df['Species'] = le.fit_transform(df['Species'])
        2 df.head()
```

```
Out[15]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [ ]: 1
```

Model Training

```
In [16]: 1 from sklearn.model_selection import train_test_split
        2
```

```
In [17]: 1 x = df.drop(columns = ['Species'])
        2
        3 x.shape
        4
```

```
Out[17]: (150, 4)
```

```
In [18]: 1 y = df['Species']
        2
        3 y.shape
```

```
In [19]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [20]: 1 # use Logistic regression
        2
        3 from sklearn.linear_model import LogisticRegression
        4 model = LogisticRegression()
```

```
In [21]: 1 # model training
        2 model.fit(x_train,y_train)
        3
```

```
Out[21]: ▾ LogisticRegression
          LogisticRegression()
```

```
In [22]: 1
        2 print("Accuracy:", model.score(x_test,y_test)*100)
```

```
Accuracy: 95.55555555555556
```

```
In [23]: 1 from sklearn.ensemble import RandomForestClassifier
        2
```

```
In [24]: 1 # create an instance of the random forest classifier
        2 clf = RandomForestClassifier(n_estimators=100)
        3
        4
```

```

In [25]: 1 # train the classifier on the training data
          2 clf.fit(x_train, y_train)
          3
          4 # predict on the test set
          5 pred = clf.predict(x_test)
          6
          7 # calculate accuracy
          8 accuracy = accuracy_score(pred, y_test)
          9 print(f"Accuracy: {accuracy}") # Accuracy: 0.91
         10

```

Accuracy: 0.9555555555555556

```

In [26]: 1 model.fit(x_train, y_train)

```

```

Out[26]: LogisticRegression
LogisticRegression()

```

```

In [27]: 1 print("Accuracy:", model.score(x_test, y_test)*100)

```

Accuracy: 95.55555555555556

```

In [28]: 1 # use DecisionTree Classifier
          2
          3 from sklearn.tree import DecisionTreeClassifier
          4 model = DecisionTreeClassifier()

```

```

In [29]: 1 model.fit(x_train, y_train)
          2

```

```

Out[29]: DecisionTreeClassifier
DecisionTreeClassifier()

```

```

In [30]: 1 print("Accuracy:", model.score(x_test, y_test)*100)

```

Accuracy: 95.55555555555556

model.py:-

```

import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import joblib

# random seed
seed = 42

# Read original dataset
iris_df = pd.read_csv("iris.csv")
iris_df.sample(frac=1, random_state=seed)

# selecting features and target data
X = iris_df[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
y = iris_df[['Species']]

# split data into train and test sets
# 70% training and 30% test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=seed, stratify=y)

# create an instance of the random forest classifier
clf = RandomForestClassifier(n_estimators=100)

```

```

# train the classifier on the training data
clf.fit(X_train, y_train)

# predict on the test set
y_pred = clf.predict(X_test)

# calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}") # Accuracy: 0.91

# save the model to disk
joblib.dump(clf, "rf_model.sav")

```

prediction.py:-

```

import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import joblib

# random seed
seed = 42

# Read original dataset
iris_df = pd.read_csv("iris.csv")
iris_df.sample(frac=1, random_state=seed)

# selecting features and target data
X = iris_df[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
y = iris_df[['Species']]

# split data into train and test sets
# 70% training and 30% test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=seed, stratify=y)

# create an instance of the random forest classifier
clf = RandomForestClassifier(n_estimators=100)

# train the classifier on the training data
clf.fit(X_train, y_train)

# predict on the test set
y_pred = clf.predict(X_test)

```

```
# calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}") # Accuracy: 0.91

# save the model to disk
joblib.dump(clf, "rf_model.sav")
```

app.py:-

```
import joblib


def predict(data):
    clf = joblib.load("rf_model.sav")
    return clf.predict(data)
```

❖ Source file:-

```
EXPLORER  ...  at Started  app.py  rf_model.sav  prediction.py  iris_flowers_classification (1).ipynb  model.py  >  ...

IRIS_PROJECT
├── __pycache__
├── .idea
├── .ipynb_checkpoints
├── app.py
├── iris_flowers_classification...
├── Iris.csv
├── model.py
├── prediction.py
└── rf_model.sav

model.py > ...
1  import pandas as pd
2  from sklearn.ensemble import RandomForestClassifier
3  from sklearn.metrics import accuracy_score
4  from sklearn.model_selection import train_test_split
5  import joblib
6  |
7  # random seed
8  seed = 42
9  |
10 # Read original dataset
11 iris_df = pd.read_csv("iris.csv")
12 iris_df.sample(frac=1, random_state=seed)
13 |
14 # selecting features and target data
15 X = iris_df[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
16 y = iris_df[['Species']]
17 |
18 # split data into train and test sets
19 # 70% training and 30% test
20 X_train, X_test, y_train, y_test = train_test_split(
21     X, y, test_size=0.3, random_state=seed, stratify=y)
22 |
23 # create an instance of the random forest classifier
24 clf = RandomForestClassifier(n_estimators=100)
25 |
26 # train the classifier on the training data
27 clf.fit(X_train, y_train)
28 |
29 # predict on the test set
30 y_pred = clf.predict(X_test)
31 |
32 # calculate accuracy
33 accuracy = accuracy_score(y_test, y_pred)
34 print(f"Accuracy: {accuracy}") # Accuracy: 0.91
35 |
36 # save the model to disk
37 joblib.dump(clf, "rf_model.sav")
```



Classifying Iris Flowers

Toy model to play to classify iris flowers into (setosa, versicolor, virginica) based on their sepal/petal and length/width.

Plant Features

Sepal characteristics

Sepal length (cm)

0.50 8.00

Sepal width (cm)

0.50 4.40

Petal characteristics

Petal length (cm)

0.50 7.00

Petal width (cm)

0.10 2.50

Iris-setosa

❖ Conclusion:-

The iris flower classification model has demonstrated a commendable performance with an accuracy score of 91.11%. This high level of accuracy indicates the effectiveness of the model in distinguishing between different species of iris flowers. The reliable classification results suggest that the model has successfully learned and generalized patterns from the input data, making it a robust tool for identifying iris species based on given features.

The accuracy score of 91.11% reflects the model's ability to make accurate predictions, thereby showcasing its potential for practical applications such as species identification in botanical research or horticulture. However, it is crucial to note that while the model has performed well, continuous evaluation and refinement may be necessary to ensure its adaptability to diverse datasets and real-world scenarios.